

1. What four main types of actions involve database? briefly discuss each?

The four main types of actions that involve a database are:

1. **Data Retrieval:** Data retrieval involves accessing data from a database. This action is typically performed when a user wants to view or analyze data that has already been stored in the database. Retrieval can be done through various methods such as SQL queries, stored procedures, and views.

Example Query: `SELECT * FROM customers WHERE age > 30;`

This query retrieves all customer records from the "customers" table where the age is greater than 30.

2. **Data Insertion:** Data insertion involves adding new data to a database. This action is typically performed when a user wants to store new information that was not previously stored in the database. Insertion can be done through SQL statements or other programming methods.

Example Query: `INSERT INTO employees (name, age, job_title) VALUES ('John Doe', 35, 'Manager');`

This query inserts a new employee record into the "employees" table with the name "John Doe", age "35", and job title "Manager".

3. **Data Modification:** Data modification involves changing existing data in a database. This action is typically performed when a user wants to update or edit information that was previously stored in the database. Modification can be done through SQL statements or other programming methods.

Example Query: `UPDATE products SET price = 10 WHERE product_id = 123;`

This query updates the price of a product with the product ID of 123 in the "products" table to \$10.

4. **Data Deletion:** Data deletion involves removing data from a database. This action is typically performed when a user wants to remove information that was previously stored in the database. Deletion can be done through SQL statements or other programming methods.

Example Query: `DELETE FROM orders WHERE order_id = 456;`

This query deletes an order record with an order ID of 456 from the "orders" table.

Overall, these four actions are the fundamental operations that can be performed on a database. Retrieval, insertion, modification, and deletion are essential for managing and maintaining the data stored in a database, ensuring that it remains accurate, up-to-date, and accessible to users who need it.

2. explain the query processing stages.

Query processing is the process of executing a database query, which involves several stages. These stages can vary depending on the type of database system and the specific query language used, but the following are common stages in most database systems:

1. **Parsing:** The first stage of query processing is parsing, which involves analyzing the query syntax and semantics to determine its structure and meaning. During parsing, the query is checked for syntax errors and translated into an internal representation that the database system can understand.
2. **Optimization:** After parsing, the query is optimized to find the most efficient way to execute it. Optimization involves evaluating different execution plans and choosing the one that requires the least amount of resources, such as CPU time or disk I/O.
3. **Semantic analysis:** This stage involves analyzing the query to ensure that it is semantically correct. The database system checks that the tables and columns referenced in the query exist and that the query doesn't violate any integrity constraints or security policies.

4. **Execution:** The final stage of query processing is execution, where the query is actually run against the database. During execution, the database system retrieves data from the storage subsystem, applies any necessary transformations or calculations to the data, and returns the results to the user.
5. **Post-processing:** After the execution of the query, the database may perform additional processing steps on the result set, such as sorting, filtering, or aggregating data. This stage is known as post-processing.

Each stage of query processing is critical for efficient and accurate database queries, and different database systems may optimize these stages in different ways.

3. Define the following terms: data, database, DBMS, database system, Userview, DBA, end user, meta-data and transaction -processing.

1. **Data:** Data refers to a collection of facts, figures, or statistics that are organized in a meaningful way. Data can be stored in a variety of forms, including text, images, audio, and video.
2. **Database:** A database is a structured collection of data that is stored in a computer system. Databases are designed to store, organize, and retrieve large amounts of information efficiently.
3. **DBMS:** DBMS stands for Database Management System, which is a software system that enables users to create, modify, and manage databases. A DBMS provides tools for data storage, retrieval, and organization, as well as security and backup features.
4. **Database system:** A database system is a collection of software, hardware, and data that is designed to manage a database. A database system includes a DBMS, as well as other components such as storage devices, backup systems, and networking equipment.
5. **Userview:** A userview is a subset of a database that is designed for a specific user or group of users. A userview can be customized to include only the data that is relevant to the user's needs.
6. **DBA:** DBA stands for Database Administrator, which is a person responsible for managing a database system. DBAs are responsible for ensuring the security, performance, and availability of the database.
7. **End user:** An end user is a person who interacts with a database system to access or modify data. End users may include customers, employees, or other stakeholders who need to access the data stored in the database.
8. **Meta-data:** Metadata is data that describes other data. In a database system, metadata includes information about the structure of the database, such as the names of tables and columns, as well as information about the data itself, such as its format and data type.
9. **Transaction processing:** Transaction processing is a technique used by database systems to ensure that multiple operations are completed successfully or not at all. Transactions are used to group related operations, such as inserting or updating data, into a single unit of work that can be committed or rolled back as a whole.

4. What are the responsibilities of the DBA and the database designers?

The DBA (Database Administrator) and database designers have distinct roles and responsibilities in managing a database system. Here are some of their primary responsibilities:

DBA Responsibilities:

1. **Database Management:** The DBA is responsible for managing the overall database system, including installation, configuration, and maintenance.
2. **Data Security:** The DBA ensures that the database is secure by establishing access controls, monitoring activity, and implementing security measures such as encryption.
3. **Backup and Recovery:** The DBA is responsible for ensuring that the database can be recovered in case of a failure or disaster by setting up backup and recovery processes.
4. **Performance Tuning:** The DBA ensures that the database system is running efficiently by monitoring performance, identifying bottlenecks, and making necessary adjustments.
5. **User Support:** The DBA provides support to users who are experiencing issues with the database system, such as connectivity or access issues.

Database Designer Responsibilities:

1. **Data Modeling:** The database designer is responsible for creating a data model that describes the structure of the database, including tables, fields, and relationships.
2. **Normalization:** The database designer ensures that the data is normalized to eliminate redundancy and improve data consistency.
3. **Data Integrity:** The database designer enforces data integrity by setting up primary keys, foreign keys, and constraints.
4. **Performance Optimization:** The database designer works to optimize the database's performance by creating efficient data models and recommending indexing strategies.
5. **Database Documentation:** The database designer maintains documentation of the database schema and any changes made to it.

Overall, the DBA and database designers work together to ensure that the database system is secure, efficient, and effective in meeting the needs of the organization.

5.What are ACID properties? what is deadlock?

ACID is an acronym that stands for Atomicity, Consistency, Isolation, and Durability. It represents a set of properties that ensure reliable transaction processing in a database system. Here is a brief explanation of each property:

1. **Atomicity:** This property ensures that a transaction is treated as a single, indivisible unit of work, either fully completed or not at all. If a transaction fails at any point, all changes made by the transaction are rolled back to their original state.
2. **Consistency:** This property ensures that a transaction brings the database from one consistent state to another. In other words, the transaction must obey all rules and constraints of the database schema and leave the database in a valid state.
3. **Isolation:** This property ensures that each transaction is executed independently of other transactions. Each transaction sees a consistent snapshot of the database, regardless of the state of other transactions.
4. **Durability:** This property ensures that once a transaction is committed, its changes are permanent and will survive any subsequent failures, such as power outages or system crashes.

Deadlock is a situation that can occur in a multi-user database system when two or more transactions are blocked because each is waiting for the other to release a resource that they need to continue executing. In other words, each transaction is waiting for the other to complete before it can proceed, resulting in a circular dependency. This can cause the database system to become unresponsive and can lead to a system crash. To avoid deadlocks, database systems use techniques such as locking and timeout mechanisms to ensure that transactions do not wait indefinitely for resources.

6. Explain Locking properties in detail?

Locking is a mechanism used in database management systems (DBMS) to provide concurrency control and prevent conflicts between transactions accessing the same data concurrently. Locking allows multiple transactions to access the same data without interfering with each other, ensuring that each transaction sees a consistent view of the data.

There are several properties of locking in DBMS:

1. **Granularity:** This refers to the level at which locks are applied. There are two types of granularities: coarse-grained locking and fine-grained locking. In coarse-grained locking, entire tables or large portions of tables are locked, whereas in fine-grained locking, individual rows or small subsets of data are locked.
2. **Mode:** This refers to the type of lock that can be acquired on a data item. There are two types of lock modes: shared locks and exclusive locks. Shared locks allow multiple transactions to read the same data concurrently but prevent any transaction from modifying the data until the lock is released. Exclusive locks prevent any other transaction from accessing the data, either for reading or writing, until the lock is released.

3. **Duration:** This refers to the length of time for which a lock is held. Locks can be held for the duration of a transaction, until the end of a statement, or for a shorter period of time, such as a few milliseconds.
4. **Deadlock:** A deadlock occurs when two or more transactions are waiting for each other to release locks, resulting in a situation where none of the transactions can proceed. DBMS uses various techniques to prevent or resolve deadlocks, such as timeout mechanisms or killing one of the transactions involved in the deadlock.
5. **Lock compatibility:** Lock compatibility refers to the ability of multiple transactions to acquire locks on the same data item. In general, exclusive locks are not compatible with other exclusive or shared locks, while shared locks are compatible with other shared locks but not with exclusive locks.
6. **Lock escalation:** Lock escalation is a mechanism by which a DBMS can convert many fine-grained locks into a smaller number of coarse-grained locks. This can reduce the overhead of managing many locks and improve performance, but can also increase the risk of contention between transactions.

7. Define the following terms:

- a. **Data model**
- b. **Database schema**
- c. **Data independence**
- d. **DDL, DML**

a. **Data model:** A data model is a conceptual representation of data that defines the structure, relationships, and constraints of data in a system. It provides a high-level view of the data and how it is organized, allowing developers and stakeholders to better understand the data and how it is used.

b. **Database schema:** A database schema is a description of the structure and organization of a database. It defines the tables, columns, relationships, constraints, and other elements that make up the database. The schema acts as a blueprint for the database, ensuring that all data is stored in a consistent and organized manner.

c. **Data independence:** Data independence refers to the ability to make changes to the database structure without affecting the applications that use it. There are two types of data independence: physical and logical. Physical data independence means that changes to the physical storage of data, such as moving data to a different disk or changing file formats, do not affect the application. Logical data independence means that changes to the logical structure of the data, such as adding or removing tables or columns, do not affect the application.

d. **DDL, DML:** DDL stands for Data Definition Language, and DML stands for Data Manipulation Language. DDL is used to define the structure of a database, including creating tables, defining columns, and setting constraints. DML is used to manipulate the data within the database, including inserting, updating, and deleting records. Examples of DDL statements include CREATE TABLE and ALTER TABLE, while examples of DML statements include SELECT, INSERT, UPDATE, and DELETE.

8. What is the difference between the two-tier and three-tier client/server architectures?

Tier-2 Architecture

The 2-tier Architecture is based on a client-server machine.

In this type of architecture, the applications on client-side interact directly with the database present at the server-side.

This interaction between client and server uses Application Program Interface like ODBC and JDBC.


- ODBC – Open Database Connectivity
- JDBC – Java Database Connectivity

When there are a large number of users at client side to access the database, this architecture gives a poor performance.

The server side is responsible for delivering the functionalities like query processing and management of transactions.

For example – Oracle, Sybase, Microsoft SQL Server etc.

The Tier-2 architecture of DBMS is diagrammatically represented as follows –



Tier-3 Architecture


The 3-tier architecture contains one more layer between the client and the server.

- In this architecture, there is no direct communication between client and server.
- Mainly, the 3-tier is used for large applications on the web.
- The features of 3-tier architecture are data backup, recovery, security, and concurrency control.

Layers

The 3-tier architecture consists of the three layers as follows –

- **Presentation layer** – This layer is also called the client layer. The front-end layer consists of a user interface. The main purpose is to communicate with the application layer.
- **Application layer** – This layer is also called the business logic layer. It acts as a middle layer between the client and the database server which are used to exchange partially processed data.
- **Database layer** – In this layer the data or information is stored. This layer performs operations like insert, update and delete to connect with the database.



9. Discuss insertion, deletion, and modification anomalies. Why are they considered bad? Illustrate with examples.

Insertion, deletion, and modification anomalies are data inconsistencies that occur when data is added, deleted, or modified in a database. These anomalies are considered bad because they can lead to incorrect, incomplete, or inconsistent data in the database. This can result in incorrect query results, wasted storage space, and difficulties in maintaining data integrity.

1. **Insertion anomaly:** An insertion anomaly occurs when it is not possible to add data to a database without creating other missing or incomplete data. For example, consider a database that stores information about students and courses. If a new student is enrolled in a course, but that student has not yet taken any exams, the student's record cannot be added to the database until the exam results are available. This creates an insertion anomaly because the database cannot handle a situation where a student is enrolled in a course but has not taken any exams.
2. **Deletion anomaly:** A deletion anomaly occurs when deleting data from the database results in the loss of other important data. For example, consider a database that stores information about employees and departments. If an employee is deleted from the database, and that employee is the only employee in a particular department, then all the information about that department is lost. This creates a deletion anomaly because the database cannot handle a situation where deleting one record results in the loss of other important data.
3. **Modification Anomaly:**

A modification anomaly occurs when a user tries to update data in a table, but doing so causes unintended changes to other rows in the table that should not be affected.

For example, suppose a company has a table that contains information about its employees, including their department and salary. If a user wants to update the salary of an employee in a particular department, but accidentally updates the salary of all employees in that department, this is a modification anomaly.

10. Why should NULL in a relation be avoided as much as possible ? Discuss the problem of [unauthentic/fake] spurious tuples and how we may prevent it

In database theory, NULL is a special marker used to indicate that a data value does not exist in the database. While NULL values can sometimes be useful, it is generally recommended to avoid using them as much as possible in a relation. Here are some reasons why:

1. **NULL values can lead to unexpected results in queries:** Since NULL values are treated as unknown, they can cause unexpected results when used in conjunction with other values in queries. For example, if you try to compare a NULL value to a numeric value, the result will be NULL, which may not be what you expect.
2. **NULL values can make it difficult to enforce constraints:** Constraints are used to ensure that the data in a database meets certain rules or requirements. When NULL values are allowed, it can make it more difficult to enforce these constraints, as it is unclear what the value should be.
3. **NULL values can lead to unnecessary complexity:** When NULL values are used, it can make the database schema more complex, as additional logic may be required to handle the NULL values correctly.

One problem that can arise as a result of using NULL values is the creation of spurious tuples, also known as phantom tuples or fake tuples. Spurious tuples are tuples that do not represent any real-world entity, but are created as a result of NULL values being present in the relation. For example, consider a relation that represents orders and has the following attributes:

(OrderID, CustomerID, OrderDate, ShipDate)

If a NULL value is allowed for the ShipDate attribute, it is possible to create a spurious tuple that represents an order that has not yet been shipped. This can cause problems when querying the database, as the spurious tuple may be included in the results, leading to incorrect information being presented to users.

To prevent spurious tuples from occurring, it is important to avoid using NULL values as much as possible. One way to do this is to use default values instead of NULL values when inserting new records into the database. Another approach is to use a separate relation to represent the absence of a value, rather than using NULL values in the main relation. Additionally, it may be helpful to enforce constraints that prevent NULL values from being used in certain situations, or to use a database management system that automatically handles NULL values in a consistent manner.

11. What is the use of Rollback Statement and how to manage it while data recovery?

12. A rollback statement is used in database management systems to undo or reverse a transaction that has not yet been committed to the database. When a transaction is started, it may consist of one or more individual database operations, such as inserts, updates, or deletions. If an error occurs during the transaction or if it is canceled before it is committed, the rollback statement is used to undo all the changes that were made as part of the transaction.
13. Managing rollback statements is an important aspect of data recovery in a database management system. In the event of a system failure or data corruption, it may be necessary to use a backup copy of the database to restore the system to a previous state. If transactions were in progress at the time of the failure, it may be necessary to use the rollback statement to undo those transactions before restoring the backup.
14. To manage rollback statements during data recovery, it is important to ensure that the database management system is configured to maintain transaction logs. These logs record all database operations that are part of each transaction, along with information about the transaction itself, such as the user who initiated it and the time it started. In the event of a failure or corruption, these logs can be used to identify and roll back any transactions that were in progress at the time of the failure.
15. To use the rollback statement to recover data, you would typically need to first restore a backup copy of the database to a new location. Then, you would use the transaction logs to identify any transactions that were in progress at the time of the failure and use the rollback statement to undo those transactions. Once the transactions have been rolled back, the restored database can be brought online and used to resume normal operations.

12. How to manage users by assigning different privileges?

How to manage users by assigning different privileges? in dbms

In a database management system (DBMS), the process of managing users and assigning different privileges typically involves the following steps:

1. Create user accounts: First, you need to create a user account for each individual user or group of users. User accounts typically include a username and password that the user will use to access the system.
2. Define roles: Next, you need to define different roles for the users in your system. Roles typically represent a set of privileges that are assigned to a particular user or group of users. For example, you may have a "read-only" role that only allows users to view data, and an "admin" role that allows users to perform administrative tasks such as creating tables and managing users.
3. Assign privileges: Once you have defined the roles, you can assign privileges to each role. This typically involves granting different levels of access to different parts of the database, such as tables, views, or stored procedures. For example, you may grant the "read-only" role access to certain tables but not others, and grant the "admin" role full access to all parts of the database.
4. Assign roles to users: Finally, you can assign the appropriate roles to each user. This typically involves associating a user account with one or more roles. For example, you may assign the "read-only" role to regular users, and the "admin" role to system administrators.

Most DBMSs have specific tools and commands for managing users and assigning privileges. For example, in MySQL, you can use the GRANT and REVOKE commands to grant and revoke privileges, and the CREATE USER and DROP USER commands to create and delete user accounts. In Microsoft SQL Server, you can use the CREATE LOGIN and CREATE USER commands to create user accounts and assign permissions using the GRANT and DENY commands. Consult the documentation for your particular DBMS for more information on managing users and assigning privileges.

13. What is meant by the concurrent execution of database transactions in a multiuser system? Discuss why concurrency control is needed and give informal example.

Concurrent execution of database transactions in a multiuser system refers to the ability of multiple users or applications to access and manipulate the same data simultaneously in a database management system. In other words, concurrent execution means that two or more transactions can execute at the same time on the same database.

Concurrency control is needed in multiuser systems to ensure that transactions are executed in a correct and consistent manner, and that the integrity of the database is maintained. Without concurrency control, transactions could interfere with each other, resulting in incorrect or inconsistent data.

For example, suppose that a banking system allows multiple users to access their accounts simultaneously. A user wants to transfer \$100 from their account to another user's account. At the same time, another user wants to deposit \$500 into their account. If these transactions are executed concurrently without proper concurrency control, the system may become inconsistent, resulting in an incorrect balance for one or more accounts. Concurrency control mechanisms can prevent this by ensuring that only one transaction is allowed to modify a particular account at a time, thereby guaranteeing consistency and accuracy.

14. Discuss the actions taken by the read_item and write_item operations on a database.

The **read_item** and **write_item** operations are two fundamental operations used in databases to retrieve data from and store data into a database, respectively. Here's how each of them works:

1. **read_item**:

The **read_item** operation is used to retrieve data from a database. When a **read_item** operation is performed, the database first checks if the requested item exists in the database. If the item exists, the database retrieves the item and returns it to the caller. If the item does not exist, the database returns an error indicating that the item was not found.

The **read_item** operation can also be used to read a specific attribute or set of attributes of an item. In this case, the database only returns the requested attributes of the item rather than returning the entire item.

2. **write_item**:

The **write_item** operation is used to store data into a database. When a **write_item** operation is performed, the database first checks if the item already exists in the database. If the item does not exist, the database creates a new item with the specified data and adds it to the database. If the item already exists, the database updates the item with the new data.

The **write_item** operation can also be used to add or remove attributes from an item. In this case, the database updates the item with the new attribute or removes the specified attribute from the item.

In both **read_item** and **write_item** operations, the database may perform various internal optimizations and checks to ensure data consistency and reliability. For example, the database may use locking mechanisms to prevent multiple users from modifying the same item simultaneously or use indexes to speed up queries. The exact implementation details depend on the specific database management system being used.

15. What is the system log used for? What are the typical kinds of records in a system log?

The system log, also known as the transaction log or the redo log, is a critical component of database management systems (DBMS). It is used to record all the transactions that occur within the database system, providing a detailed history of changes made to the database.

The system log is used for several purposes, including:

1. **Recovery:** The system log enables the DBMS to recover from a system failure or a transaction failure by allowing the system to roll back or roll forward to a known, consistent state.
2. **Replication:** The system log is used in replication scenarios, where a copy of the database is maintained on multiple servers. The system log enables the replica servers to keep up with changes made to the primary database.
3. **Auditing:** The system log provides a record of all the transactions that have occurred within the database system, making it useful for auditing and compliance purposes.

The typical kinds of records found in a system log include:

1. **Transaction ID:** A unique identifier for each transaction.
2. **Timestamp:** The date and time when the transaction occurred.
3. **Operation:** The type of operation performed, such as insert, update, or delete.
4. **Object ID:** The identifier of the database object (e.g., table, index) affected by the transaction.
5. **Data:** The data that was inserted, updated, or deleted.
6. **Status:** The status of the transaction, such as committed, aborted, or in progress.

By analyzing the system log, DBAs can identify performance bottlenecks, troubleshoot issues, and track down data inconsistencies.

16 what is a serial schedule ? what is a serializable schedule? Why is a serial schedule considered correct ? why a serializable schedule is considered correct ?

A serial schedule is a type of schedule in database management systems where the transactions are executed one after the other, with no overlapping execution. In other words, a serial schedule is a schedule in which transactions are executed in a sequential and non-concurrent manner.

A serializable schedule is a type of schedule in database management systems where concurrent transactions are executed in such a way that the final result is equivalent to a serial schedule. In other words, a serializable schedule is a schedule in which transactions are executed concurrently but produce the same result as if they were executed sequentially.

A serial schedule is considered correct because it guarantees that the transactions will be executed in a consistent and predictable manner, with no conflicts or interference between them. A serializable schedule is also considered correct because it ensures that concurrent transactions will produce the same result as if they were executed sequentially, even though they are executed concurrently.

The difference between a serial schedule and a serializable schedule is that a serial schedule executes transactions one after the other, while a serializable schedule allows transactions to be executed concurrently as long as the final result is equivalent to a serial schedule. In other words, a serializable schedule is a more efficient way of executing transactions because it allows for concurrency while still maintaining consistency and correctness.

17. Define first, second, and third normal forms when only primary keys are considered.

First, second, and third normal forms are related to the elimination of data redundancy and dependencies in a relational database. These normal forms are defined based on the concept of functional dependencies between attributes in a table.

When only primary keys are considered, the definitions of the normal forms are as follows:

1. First Normal Form (1NF): A table is said to be in 1NF if all its attributes are atomic and do not contain repeating groups. In other words, each column in a 1NF table contains only one value, and each row is unique.
2. Second Normal Form (2NF): A table is in 2NF if it is already in 1NF and every non-primary key column is fully functionally dependent on the primary key. This means that each non-primary key column must depend on the entire primary key and not just a part of it.
3. Third Normal Form (3NF): A table is in 3NF if it is already in 2NF and every non-primary key column is independent of every other non-primary key column. This means that there should be no transitive functional dependencies between non-primary key columns.

In summary, when only primary keys are considered, a table is said to be in 1NF if it has no repeating groups and each column contains only one value, in 2NF if each non-primary key column is fully dependent on the entire primary key, and in 3NF if there are no transitive dependencies between non-primary key columns.

18. what are transaction commit points, and why are he important?

Transaction commit points are specific points in a database transaction where the changes made to the database are permanently saved or "committed".

Transaction commit points are points in a database transaction where the changes made to the database are permanently saved and become visible to other transactions. In other words, a transaction commit point is the point at which a transaction is completed, and its changes are made permanent in the database.

Commit points are important for several reasons:

1. Data consistency: By committing changes at specific points in a transaction, data consistency can be maintained. If a transaction fails and must be rolled back, only the changes made since the last commit point will be lost, not the entire transaction.
2. Database recovery: Commit points are important for database recovery in the event of a system failure or crash. By committing changes at specific points, the database can be restored to a consistent state by undoing any changes made after the last commit point.
3. Concurrent access: When multiple users are accessing the same database simultaneously, commit points ensure that other users can see changes made by a transaction only after the transaction is committed. This ensures that concurrent users do not see intermediate or inconsistent states of the database.

Overall, commit points are important for ensuring data consistency and integrity, as well as for database recovery and concurrent access. They provide a way to break up a transaction into smaller, more manageable units and provide a clear point at which to commit changes to the database.

19 what are the different types of constraints used in SQL?

Constraints are used in SQL to define rules and restrictions that must be followed while inserting, updating, or deleting data from a database. There are several types of constraints used in SQL, including:

1. NOT NULL constraint: This constraint ensures that a column cannot contain a NULL value.
2. UNIQUE constraint: This constraint ensures that each value in a column is unique and cannot be repeated.
3. PRIMARY KEY constraint: This constraint ensures that each row in a table is uniquely identified by a column or set of columns, and cannot be duplicated.
4. FOREIGN KEY constraint: This constraint is used to ensure referential integrity between two tables. It ensures that a value in one table's column (referred to as the foreign key) matches a value in another table's column (referred to as the primary key).
5. CHECK constraint: This constraint ensures that the values in a column meet a specific condition or set of conditions.
6. DEFAULT constraint: This constraint provides a default value for a column if no value is specified during an INSERT operation.

7. **INDEX constraint:** This constraint is used to create an index on one or more columns of a table to improve query performance.

These constraints can be used individually or in combination to ensure data integrity and consistency in a database.

20.what is a schedule? Define the concepts of recoverable, cascade less, and strict schedules and compare them in terms of their recoverability.

A schedule is a sequence of operations or transactions that are executed by a database management system (DBMS). In the context of databases, a schedule refers to a list of database operations that are executed in a specific order.

Recoverable schedule: A recoverable schedule is a schedule in which if a transaction fails and is rolled back, all the changes made by that transaction are undone. In other words, the system can recover from a failure without losing any committed data.

Cascadeless schedule: A cascadeless schedule is a schedule in which no transaction reads data that has been updated but not yet committed by another transaction. This ensures that the system does not get into a state where one transaction depends on the uncommitted data of another transaction.

Strict schedule: A strict schedule is a schedule in which no transaction reads data that has been written but not yet committed by another transaction, and no transaction writes data that has been read by another transaction but not yet committed. This ensures that the system is always in a consistent state.

In terms of recoverability, a strict schedule is the most recoverable because it ensures that the system is always in a consistent state. A cascadeless schedule is also recoverable because it ensures that the system does not get into a state where one transaction depends on the uncommitted data of another transaction. However, a recoverable schedule may not be cascadeless or strict. In a recoverable schedule, a transaction may read uncommitted data from another transaction, which could lead to inconsistencies in the system. Therefore, recoverable schedules may require additional measures to ensure consistency and recoverability.